

1 02805 Social graphs and interactions: Project 2 Assignment B

3 **Ioannis Vlasakoudis s232755, Max Heiberg Bestle s194574, Jonatan Rasmussen s183649**

4 December 11, 2024

5 In this paper, we analyze the 20 Newsgroups dataset to understand how online users interacted
6 across different topics. We construct a network where nodes represent users who posted
7 across the newsgroups, with edges connecting users based on their interactions. We then
8 apply community detection algorithms to uncover natural groupings of users.

9 Social networks | Text analysis | Sentiment Analysis | TF-IDF | Data Science

10 In the early days of the internet, online communities formed through newsgroups
11 where people shared opinions, discussed various topics, and engaged in conversations
12 through email-like posts.

13 These newsgroups represented one of the first large-scale digital social networks,
14 making them an interesting subject for both network and text analysis. In this
15 paper, we analyze the 20 Newsgroups dataset to understand how these early online
16 communities were structured and how users interacted across different topics.(1).

24 Research questions

25 Our analysis combines network science with natural language processing to
26 investigate three main questions: How did users form communities across different
27 newsgroups, and are they aligning with the predefined newsgroup categories? What
28 were the main topics of discussion within and across these communities? And how
29 did the sentiment of discussions vary across different topics and communities?

31 Methodology

32 To answer these questions, we constructed a network where nodes represent users
33 who posted across the newsgroups, with edges connecting users who interacted
34 with each other. We then applied the Louvain method to uncover the structural
35 communities of the users. These communities are analyzed using both topic modeling
36 techniques (Latent Dirichlet Allocation, more on this later) and sentiment analysis
37 to understand the nature of discussions within each community. For sentiment
38 analysis, we use the LabMT wordlist(2) for assigning a 'happiness rank' to common
39 English words, which allows us compare the positivity/negativity across forums.
40 *Link to our notebook: <https://github.com/maxx1559/socialgraphs-final/tree/main>*

42 Data overview and source

43 As previously mentioned, our data has been sourced from the 20 Newsgroups
44 dataset. The version of the dataset that we are using has been downloaded
45 from Kaggle.com, see the references section.(3) In terms of dataset content, each
46 newsgroup is discussing a specific subject, such as religion, politics, or sport. An
47 example of what posts look like can be seen in figure 1.

48 Each newsgroup submission has the following attributes: document id, the
49 newsgroup, the author, and subject title. However, the dataset is a big text file
50 containing 1000+ posts, so we decided to do some pre-processing and turn it into
51 JSON files. When splitting the raw data into posts, we ended up with a total of
52 18,773 posts. The posts are distributed approximately evenly across the newsgroups,
53 with most containing around 1,000 samples, see figure 2. The average length of
54 each post is 193 words, which is quite long by modern social media standards. This
55 means that we have almost 4 million words in our data set!

56 **Data partitioning.** The 20 Newsgroups website suggests grouping the newsgroups
57 into six broader subject areas: Computers, Buying & Selling, Recreational, Politics,
58 Science, and Religion, see figure 3. Generally speaking, it is also easier to display
59 6 categories in a plot compared to 20 categories, which is why we are using this
60 partition for parts of our work.

71 Significance Statement

72 This project work provides insights into how early internet users
73 formed social bonds and communities around shared interests, and
74 how these communities sometimes transcended the formal boundaries
75 of individual newsgroups. Understanding these patterns can help us
76 better comprehend the evolution of online social networks and community
77 formation in digital spaces.

63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124

```

125 Newsgroup: comp.os.ms-windows.misc
126 document_id: 10041
127 From: bjaastad@idt.unit.no (H)vard Bjastad)
128 Subject: HELP! Word4W Sucks !?!?!? (AUTONUM!)
129
130 As I am working on th report I hardly have any time to read News, so
131 please e-mail me. All answers will be heartly welcome ...
132
133 In advance, thanx a lot !!!!
134 -----
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186

```

Frode Rinnans v.9, 7035 Trondheim
Tlf.: 07588723
bjaastad@idt.unit.no

"And on the 8th day,

Fig. 1. An example of a newsgroup post. This is what our data looks like!

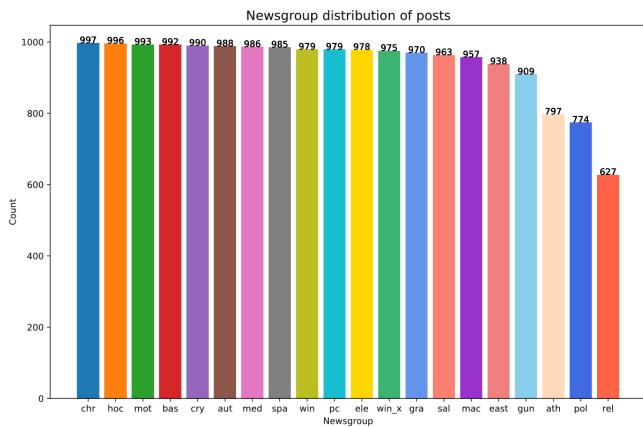


Fig. 2. Distribution of posts across newsgroups. It appears that the original dataset sampled 1000 posts from each newsgroup. However the religion.misc newsgroup stands out as an outlier, having only 627 samples. Still, this should not cause any issues for our data analysis work, as it's a very healthy sample size.

Results and key findings

This section is a presentation and visualizes of our key findings. Discussion regarding our results can be found in the discussion section on the next page of the report. For our code and additional data, check out the notebook.(4)

Individual Networks. In order to construct the large network between users in all newsgroups, we first had to analyze and prepare the individual newsgroup networks. As mentioned in Methodology, we created the networks, with nodes representing the newsgroup's users and they were connected to each other if there was an interaction between them (for example, a user mentioned the name of another user/node in their post). For each newsgroup, only the largest connected component was selected for analysis and the result was 20 networks with an average of 246 nodes and 574 edges, with the largest being the network of *Computer Graphics* with 636 nodes and 1222 edges (see Figure 4).

Unified Network and Communities. After the analysis of the individual networks, it is time to create the large unified network. For that, we had to begin by identifying the common users across these 20 networks. There are 347 users that have posted in different newsgroups and most of them between atheism and religion (32). After the combination, a unified network of 4306 nodes and 11098 edges emerged (Figure 5).

comp.graphics	rec.autos	sci.crypt	187
comp.os.ms-windows.misc	rec.motorcycles	sci.electronics	188
comp.sys.ibm.pc.hardware	rec.sport.baseball	sci.med	189
comp.sys.mac.hardware	rec.sport.hockey	sci.space	190
comp.windows.x			191
misc.forsale	talk.politics.misc	talk.religion.misc	192
	talk.politics.guns	alt.atheism	193
	talk.politics.mideast	soc.religion.christian	194

Fig. 3. The 20 newsgroups partitioned into 6 subjects.

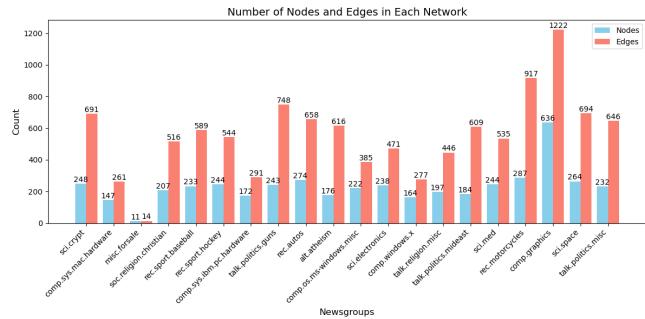


Fig. 4. Number of Nodes and Edges of Individual Networks

The next step was to apply the Louvain method and detect the communities, which could offer us valuable insights. From this method, 21 communities were derived. Afterwards, we proceeded with the analysis of these communities and calculated the distribution of users, as well as the total number of common users in each community (Figure 6). Looking at the distribution of degrees within the large network, we found that while the network did seem to follow a power-law distribution, it did not follow a scale-free regime, for either in-degrees, out-degrees or the total degrees.

TF-IDF and Most Frequent Words. Analyzing the word frequency is an additional way to deeply understand the communities and how the users in them think. Thus, we found the 10 most used words for each community (after some text processing) and for a better, comprehensive analysis, we also calculated the TF-IDF scores of the words, in order to identify which words are the most important for each community. For example, the top words for community 3 are *'gun'*, *'people'*, *'one'*, *'government'*, *'right'*, *'get'*, *'think'*, *'weapon'*, *'law'*, *'fire'*, and they can reveal a lot about the context of the posts.

Sentiment Analysis. To measure the language used across different topics, we have used sentiment analysis to quantify the positivity/negativity of each topic. During our analysis, we found that the *forsale* newsgroup had an average sentiment score of 5.566 (± 0.179), which was the highest of the sentiments, whereas the political newsgroup had the lowest average sentiment score, which was at 5.312 (± 0.162). The sentiment all the subjects can be seen in figure 7. We see that the difference is statistically significant, which means that the sentiment actually differs across topics. Perhaps it is not too surprising that for-sale posts contain positive words and language to promote the items that are being sold, whereas political discussion is more likely to spark conflicts and disagreeable language.

249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310



Fig. 5. The Unified Network

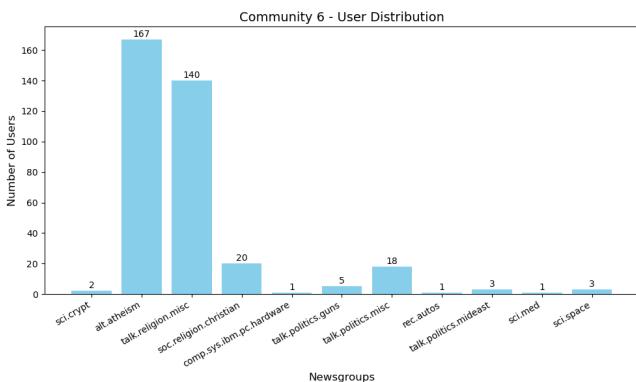


Fig. 6. User Distribution of Community 6

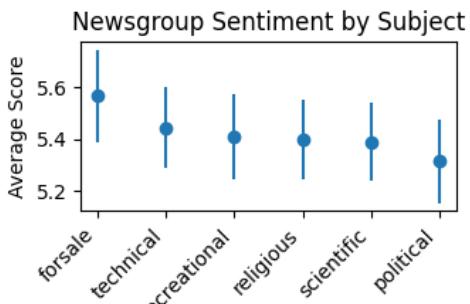


Fig. 7. Newsgroups sentiment by subject. We observe that the *forsale* newsgroup has the highest sentiments, whereas the partitioned political newsgroups have the lowest sentiment. The other four partitioned newsgroups have nearly identical sentiment, which is also somewhat surprising as they differ widely in subject.

Community Sentiments. Now that we have the sentiments for each newsgroup, we can use them as a comparison metric for the sentiment analysis of the communities. The highest sentiment is that of community 17 with 5.683 (± 0.854), which has users who have only posted in the newsgroup regarding Windows operating system. Community 3 has the lowest sentiment score of 5.350 (± 1.147) and includes users with interests in guns, politics, religion, space, and other topics.

LDA and latent topics. Latent Dirichlet Allocation (LDA) is an algorithm that can discover "hidden" topics in any text by counting which words frequently appear together. The output of the algorithm is a list of the most frequent words for each of the self-discovered n topics. For example, we can set $n=6$ and execute the LDA-algorithm on our entire dataset to discover 6 latent topics which may or may not be the same word distributions as our partitioned newsgroups. Noteworthy findings include Topic 5, where the top words are "armenian", "turkish", "israel", "israeli", "arab", "jew", "turk", "said" and "war". Without using any of the newsgroup labeling, the computer found these terms to be especially tightly linked. Now compare that to Topic 0, which is (*would*, *one*, *get*, *like*, *dont*, *time*, *good*, *know*, *im*, *could*) or Topic 4 (people, one, would, god, say, think, dont, know, u, right) which from a human standpoint seems much more random. In that context, it is remarkable that Topic 5 is a word clustering that both humans and the LDA algorithm agrees belong together.

Discussion

Here we present the key insights gained from the analysis. Starting with the individual newsgroup networks, the first thing to be noticed is that almost all networks are balanced in terms of node numbers, except for newsgroup *for sale*, which is a very small network and this is understandable, as we can assume that people in this category mainly communicated with specific users who were selling something. Thus, a big network was unlikely to be created here. Furthermore, looking at the top nodes/users by degree in each newsgroup network, we notice some interesting patterns. In car and computer related topics, the out-degree of the top users is, in most cases, significantly higher than their in-degree. This indicates that they initiate interactions more frequently than they are contacted/messaged by others. This could be happening because these users have the knowledge to give advice or solve technical problems that other users have, related to these certain topics. In general, the in/out-degree distributions tend to be either relatively balanced or the out-degree being higher, which means that the top users of many newsgroups mainly reach out to people, but it is, in the majority, a back and forth conversation.

Looking within each of the individual networks, we looked at the communities that might arise in each newsgroup. Here, most of the networks have a fairly high modularity, the lowest of them still being over 0.4, showing that communities within them do seem to appear. This could be because the communities are created by small back and forth interactions between users, meaning they often reference each other in their posts. Some users could also represent one single often referred to post, meaning their e-mail is often represented in other users' posts.

Moving to the large unified network, it is a very sparse network with a *Density* of 0.0011, which is a common characteristic among real-life social networks. Looking at the degree distribution, we can see that the degree exponents are all along the lines of $\gamma > 3$ for either in- or out-degrees, as well as total degrees. This means the network is more or less random, or at least hard to distinguish from a random network of the same size(5). This mirrors the network for emails, where it still follows a power law distribution, but is more or less random. Another characteristic that indicates the existence of a real social network, is the high level of local connectivity between the users (*Average Clustering Coefficient*: 0.4731). That is how local communities are formed.

And talking about communities, we can draw multiple conclusions from analyzing the communities derived by the Louvain method. First and foremost, we can see that the *Modularity* is very high (0.85), which indicates that the network has well-defined communities with strong structures. This was an expected result, if we think the way this network was build (combination of smaller networks through common nodes). Looking at the community user distributions, we can observe that users in them predominantly talk/post about a single topic, except for some specific communities. In these specific communities, the users talk/post about relatively similar topics. For example, in community 14, the users are talking about computer hardware, software and electronics or in community 6, they are talking about atheism, Christianity and politics. Thus, we can conclude that the communities contain users with the same interests like technology, social matters, sports etc..

Regarding the text analysis of each community, it can help us understand what the main context of the posts is, without even looking at the user distribution. Especially for simple word frequencies, understanding the topic is very easy by looking at the top 10 words of a community. On the other hand, looking at the words with the highest TF-IDF scores, does not provide as much insight into the topics and interests of a community as we initially thought. For most of the communities, we get terms that could be characterized as key words for the topic of a community, but there are also some words like names or other nonsense words that cannot help us distinguish one community from another. For example, in communities . Thus, we can use this only as a complementary analysis

Last but not least, we have the sentiment analysis. With this analysis, we can understand what is the overall "mood" of each community. To justify the results, we are going to use also the sentiment analysis of the newsgroups. It can be noticed that when the interests of a community are politics (especially Middle-Eastern) or guns, there is a negativity in the posts, while for Christianity or technical topics there is a positivity. There are also some examples, like community 19, which ranked third in overall sentiment and the context of the posts was about religion. Although religion ranks in the middle to low range in terms of sentiment, this community is ranked quite high. This indicates that some communities are also divided in terms of sentiment. By this, we mean that certain communities contain users who interacted with each other and the context of their conversation was positive, even though the general topic of their conversation may have a low sentiment score.

Conclusion

In terms of our three research questions mentioned at the start of the paper, we have reached the following conclusions: *How did users form communities across different newsgroups, and are they aligning with the predefined newsgroup categories?* While users often stayed within their primary interest areas, we found significant cross-posting between related topics. The 21 communities detected by the Louvain method showed that users tended to group around broader interest areas rather than strictly following newsgroup boundaries. *What were the main topics of discussion within and across these communities?* Our text analysis revealed both highly focused single-topic communities and broader multi-topic discussions. The LDA analysis uncovered some surprising topic clusters, showing deeper patterns in how discussions naturally organized themselves.

How did sentiment vary across different topics and communities? We found statistically significant variations in sentiment across topics, with political discussions being notably more negative than buy-and-sell discussions. Community sentiments generally reflected their primary topics, though some communities showed unique emotional characteristics independent of their subject matter.

Overall, we found that users don't necessarily tend to stay within the same topics. Many users post in different topics, and because of this, communities do seem to arise based on shared interests. Additionally, our sentiment analysis found some of the newsgroups to be more unhappy than others, with the political communities being the most unhappy, while the buy and sell and technology related communities seemingly using the most positive language.

Contributions

Max: Data pre-processing and preliminary network analysis
Ioannis: Network analysis and community analysis
Jonatan: Newsgroups text, sentiment analysis and LDA

Jupyter notebook

Link to our notebook:
<https://github.com/maxx1559/socialgraphs-final/tree/main>

References

1. Ken Lang. The 20 newsgroups data set. <http://qwone.com/~jason/20Newsgroups/>, 2007. Accessed: 2024-12-10.
2. Peter Sheridan Dodds, Kameron Decker Harris, Isabel M. Kloumann, Catherine A. Bliss, and Christopher M. Danforth. Temporal patterns of happiness and information in a global social network: Hedonometrics and twitter. <https://doi.org/10.1371/journal.pone.0026752>, 2011. Published: December 7, 2011, accessed 2024-12-10.
3. Chris Crawford. 20 newsgroups: A collection of 18,000 newsgroup documents from 20 different newsgroups. <https://www.kaggle.com/datasets/chrisc/c20-newsgroups>, 2017. Accessed: 2024-12-10, updated 7 years ago.
4. Ioannis Vlasakoudis, Max Heilberg Bestle, and Jonatan Rasmussen. Python notebook with code for this project. <https://github.com/maxx1559/socialgraphs-final/tree/main>, 2024. Accessed: 2024-12-10.
5. Albert-Lazlo Barabasi. Network science. <https://networksciencebook.com/chapter/1>, 2016. Accessed: 2024-12-10.